

Bayesian modelling

Markov chain Monte Carlo

Léo Belzile

Last compiled Tuesday Sep 29, 2026

Markov chain Monte Carlo

We consider simulating from a distribution with associated density function proportional to $p(\cdot)$

- using an algorithm that generates a Markov chain ,
- without requiring knowledge of the normalizing factor.

We use a conditional density $q(\boldsymbol{\theta} \mid \boldsymbol{\theta}^{\text{cur}})$ to generate proposals from the current value $\boldsymbol{\theta}^{\text{cur}}$.

Metropolis–Hastings algorithm

Starting from an initial value $\boldsymbol{\theta}_0$: for $t = 1, \dots, T$

1. draw a proposal value $\boldsymbol{\theta}_t^* \sim q(\boldsymbol{\theta} \mid \boldsymbol{\theta}_{t-1})$.

2. Compute the acceptance ratio

$$R = \frac{p(\boldsymbol{\theta}_t^*)}{p(\boldsymbol{\theta}_{t-1})} \frac{q(\boldsymbol{\theta}_{t-1} \mid \boldsymbol{\theta}_t^*)}{q(\boldsymbol{\theta}_t^* \mid \boldsymbol{\theta}_{t-1})}$$

3. With probability $\alpha = \min\{R, 1\}$, accept the proposal and set $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_t^*$, otherwise set the value to the previous state, $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1}$.

Theory

The Metropolis–Hastings algorithm satisfies a technical condition (detailed balance) that ensures that the Markov chain generated is reversible.

- Thus, any draw from the posterior will generate a new realization from the posterior.
- Provided the starting value has non-zero probability under the posterior, the chain will converge to the stationarity distribution (albeit perhaps slowly).

Interpretation

- If $R > 1$, the proposal has higher density and we always accept the move.
- If we reject the move, the Markov chain stays at the current value, which induces autocorrelation.
- Since the acceptance probability depends only on the density through ratios, normalizing factors of p and q cancel out.

Symmetric proposals and random walk

If the proposal is symmetric, the ratio of proposal densities is

$$q(\boldsymbol{\theta}_{t-1} \mid \boldsymbol{\theta}_t^*) / q(\boldsymbol{\theta}_t^* \mid \boldsymbol{\theta}_{t-1}) = 1.$$

Common examples include random walk proposals

$$\boldsymbol{\theta}_t^* \leftarrow \boldsymbol{\theta}_{t-1} + \tau Z,$$

where $Z \sim \text{Gauss}(0, 1)$.

Independent proposals

- If we pick instead a global proposal, we must ensure that q samples in far regions (recall rejection sampling), otherwise ...
- Good proposals include heavy tailed distribution such as Student- t with small degrees of freedom,
- centered at the maximum a posteriori $\hat{\theta}$ and
- with scale matrix proportional to $-\mathbf{H}^{-1}(\theta_t^*)$, where $\mathbf{H}(\cdot)$ is the Hessian of the (unnormalized) log posterior.

Upworthy data example

We model the Poisson rates for headlines with questions or not. Our model is

$$Y_i \sim \text{Poisson}(n_i \lambda_i), \quad (i = 1, 2)$$

$$\lambda_1 = \exp(\beta + \kappa)$$

$$\lambda_2 = \exp(\beta)$$

$$\beta \sim \text{Gauss}(\log 0.01, 1.5)$$

$$\kappa \sim \text{Gauss}(0, 1)$$

Implementation details: data and containers

In regression models, scale inputs if possible.

```
1 # Load data
2 data(upworthy_question, package = "hecbayes")
3 # Compute sufficient statistics
4 data <- upworthy_question |>
5   dplyr::group_by(question) |>
6   dplyr::summarize(ntot = sum(impressions),
7                     y = sum(clicks))
8 # Create containers for MCMC
9 niter <- 1e4L
10 chain <- matrix(0, nrow = niter, ncol = 2L)
11 colnames(chain) <- c("beta", "kappa")
12 logpost_v <- numeric(niter)
```

Implementation details: log posterior function

Perform all calculations on the log scale to avoid numerical overflow!

```
1 # Code log posterior as sum of log likelihood and log prior
2 loglik <- function(par, counts = data$y, offset = data$ntot, ...){
3   lambda <- exp(c(par[1] + log(offset[1]),
4                   par[1] + par[2] + log(offset[2])))
5   sum(dpois(x = counts, lambda = lambda, log = TRUE))
6 }
7 # Note common signature of function
8 logprior <- function(par, ...){
9   dnorm(x = par[1], mean = log(0.01), sd = 1.5, log = TRUE) +
10   dnorm(x = par[2], log = TRUE)
11 }
12 logpost <- function(par, ...){
13   loglik(par, ...) + logprior(par, ...)
14 }
```

Implementation details: proposals

Use good starting values for your Markov chains, such as maximum a posteriori.

```
1 # Compute maximum a posteriori (MAP)
2 map <- optim(
3   par = c(-4, 0.07),
4   fn = logpost,
5   control = list(fnscale = -1),
6   offset = data$ntot,
7   counts = data$,
8   hessian = TRUE)
9 # Use MAP as starting value
10 cur <- map$par
11 # Keep track of log posterior to reduce calculations
12 logpost_cur <- logpost(cur)
13 # Proposal covariance
14 cov_map <- -2*solve(map$hessian)
15 chol <- chol(cov_map)
```

Implementation details: Metropolis–Hastings

Use seed for reproducibility, do not compute posterior twice, compute log of acceptance ratio.

```
1  set.seed(80601)
2  naccept <- 0L
3  for(i in seq_len(niter)){
4    # Multivariate normal proposal - symmetric random walk
5    prop <- c(rnorm(n = 2) %*% chol + cur)
6    logpost_prop <- logpost(prop)
7    logR <- logpost_prop - logpost_cur
8    if(logR > -rexp(1)){# accept move
9      cur <- prop
10     logpost_cur <- logpost_v[i] <- logpost_prop
11     naccept <- naccept + 1L
12   }
13   logpost_v[i] <- logpost_cur
14   chain[i,] <- cur
15 }
```

Implementation details: analysis of output

```
1 # Posterior summaries
2 summary(coda::as.mcmc(chain))
```

1. Empirical mean and standard deviation for each variable, plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
beta	-4.51268	0.001697	1.697e-05	6.176e-05
kappa	0.07075	0.002033	2.033e-05	9.741e-05

2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
beta	-4.51591	-4.51385	-4.51273	-4.51154	-4.50929
kappa	0.06673	0.06933	0.07077	0.07212	0.07463

Metropolis–Hastings algorithm

Starting from an initial value $\boldsymbol{\theta}_0$:

1. draw a proposal value $\boldsymbol{\theta}_t^* \sim q(\boldsymbol{\theta} \mid \boldsymbol{\theta}_{t-1})$.
2. Compute the acceptance ratio

$$R = \frac{p(\boldsymbol{\theta}_t^*)}{p(\boldsymbol{\theta}_{t-1})} \frac{q(\boldsymbol{\theta}_{t-1} \mid \boldsymbol{\theta}_t^*)}{q(\boldsymbol{\theta}_t^* \mid \boldsymbol{\theta}_{t-1})}$$

3. With probability $\min\{R, 1\}$, accept the proposal and set $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_t^*$, otherwise set the value to the previous state, $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1}$.

Calculations

We compute the log of the acceptance ratio, $\ln R$, to avoid numerical overflow, with the log posterior difference

$$\ln \left\{ \frac{p(\boldsymbol{\theta}_t^*)}{p(\boldsymbol{\theta}_{t-1})} \right\} = \ell(\boldsymbol{\theta}_t^*) + \ln p(\boldsymbol{\theta}_t^*) - \ell(\boldsymbol{\theta}_{t-1}) - \ln p(\boldsymbol{\theta}_{t-1})$$

Compare the value of $\ln R$ (if less than zero) to $\log(U)$, where $U \sim \text{unif}(0, 1)$.

What proposal?

The *independence* Metropolis–Hastings uses a global proposal q which does not depend on the current state (typically centered at the MAP)

- This may be problematic with multimodal targets.

The Gaussian random walk takes $\boldsymbol{\theta}_t^* = \boldsymbol{\theta}_{t-1} + \sigma_p Z$, where $Z \sim \text{Gauss}(0, 1)$ and σ_p is the proposal standard deviation. *Random walks* allow us to explore the space.

Breaking down the Markov chain

We distinguish between three phases

- **burn in** period: initial draws allowing the algorithm to converge to its stationary distribution (discarded)
- **warmup** adaptation period: tuning period for the proposal std. deviation, etc. (discarded)
- **sampling** period: draws post burn in and warmup that are kept for inference

We can optionally **thin** by keeping one every k iterations from the sampling period to reduce the storage.

Burn in

We are guaranteed to reach stationarity with Metropolis–Hastings eventually, but it may take a large number of iterations...

One should discard initial draws during a **burn in** or warmup period if the chain has not reached stationarity. Ideally, use good starting value to reduce waste.

We can also use the warmup period to adapt the variance of the proposal.

Goldilock principle and proposal variance

Mixing of the chain requires just the right variance (not too small nor too large).




Figure 1: Example of traceplot with proposal variance that is too small (top), adequate (middle) and too large (bottom).

Correlograms of MCMC output




Figure 2: Correlogram for the three Markov chains.

Tuning Markov chain Monte Carlo

- Outside of starting values, the variance of the proposal has a huge impact on the asymptotic variance.
- We can adapt the variance during warmup by increasing/decreasing proposal variance (if acceptance rate is too large/small).
- We can check this via the acceptance rate (how many proposals are accepted).

Optimal acceptance rates

The following rules were derived for Gaussian targets under idealized situations.

- In 1D, rule of thumb is an acceptance rate of 0.44 is optimal, and this ratio decreases to 0.234 when $D \geq 2$ (Sherlock, 2013) for random walk Metropolis–Hastings.
- Proposals for D -variate update should have proposal variance of roughly $(2.38^2 / d) \times \Sigma$, where Σ is the posterior variance.
- For MALA (see later), we get 0.574 rather than 0.234

Block update or one parameter at a time?

As with any accept-reject, proposals become inefficient when the dimension D increase.

This is the **curse of dimensionality**.

Updating parameters in turn

- increases acceptance rate (with clever proposals),
- but also leads to more autocorrelation between parameters

Solutions for strongly correlated coefficients

- Reparametrize the model to decorrelate variables (orthogonal parametrization).
- Block updates: draw correlated parameters together
 - using the chain history to learn the correlation, if necessary

Parameter transformation

Parameters may be bounded, e.g. $\theta_i \in [a, b]$.


- We can ignore this and simply discard proposals outside of the range, by setting the log posterior at $-\infty$ outside $[a, b]$
- We can do a transformation, e.g., $\log \theta_i$ if $\theta_i > 0$ and perform a random walk on the unconstrained space: don't forget Jacobians for $q(\cdot)$!
- Another alternative is to use truncated proposals (useful with more complex algorithms like MALA)

Visual diagnostic: trace plots

Display the Markov chain sample path as a function of the number of iterations.

- Ideally, run multiple chains to see if they converge to the same mode (for multimodal behaviour).
- Markov chains should look like a fat hairy caterpillar!
- Check the [bayesplot](#) and [coda](#) R packages (trace plot, trace rank, correlograms, marginal densities, etc.)

Checking convergence with multiple chains



Four healthy parallel chains for parameters.

Trace rank plot

A **trace rank** plot compares the rank of the values of the different chain at a given iteration.

- With good mixing, the ranks should switch frequently and be distributed uniformly across integers.

Lack of stationarity




Figure 3: Traceplots of three Markov chains for the same target with different initial values for the first 500 iterations (left) and trace rank plot after discarding these (right). The latter is indicative of the speed of mixing.

Summary (1/3)

Summary (2/3)

- Good initial starting values can be obtained by computing maximum a posteriori estimates.
- Initializing multiple chains at different starting values can be used to check convergence to the stationary distribution.
- Mixing will improve if strongly correlated parameters are sampled together.

Summary (3/3)

References

- Neal, R. M. (2011). MCMC using Hamiltonian dynamics. In S. Brooks, A. Gelman, G. Jones, & X. L. Meng (Eds.), *Handbook of Markov chain Monte Carlo* (pp. 113–162). CRC Press. <https://doi.org/10.1201/b10905-5>
- Sherlock, C. (2013). Optimal scaling of the random walk Metropolis: General criteria for the 0.234 acceptance rule. *Journal of Applied Probability*, 50(1), 1–15. <https://doi.org/10.1239/jap/1363784420>